

# Migrate when necessary: toward partitioned reclaiming for soft real-time tasks

RTNS'2015

H. ZAHAF<sup>1</sup>, G. LIPARI<sup>1</sup>, L. ABENI<sup>2</sup>

<sup>1</sup> CRIStAL, Lille 1 University

<sup>2</sup> Scuola Superiore Sant'Anna

Grenoble, 04/10/2017

# Plan

---

Introduction

Our Contributions

Results and discussions

Conclusion

# Soft Real-time

---

- Undesirable to miss deadlines but not crucial
- Modern operating systems can serve this kind of applications
- Examples
  - Multimedia coding/decoding/transmission
  - Networking, cellular operations,
  - ...
- Execution time is not known a priori
- Tasks may have a **unpredictable** behavior.

## Consequence

- Non-guarantee to respect all deadlines
- Tasks with long execution times may monopolize the processor, and then increase deadline miss rate.
- → *special* scheduling techniques

# Resource Reservation and reclaiming

---

## Temporal isolation

- Ressource reservation allow to have temporal isolation:
  - To each task is assigned a **service time** “budget”  $Q_i$  in each period of time  $P$
  - Tasks running out their budgets must not condition other's tasks execution.

## Reclaiming

- What about tasks that runs less then their budgets?
  - Need a reclaiming policy to take benifit from the *unused* budget.

# Reclaiming: Singlecore and multicore

---

## Single-core

- Several algorithms exist for single-core resource reservation and reclaiming:
  - Grub, CBS, CASH, ...

## Multicore support

- Multicore techniques are: global, partitioned or semipartitioned.
- Several extension of GRUB and CBS has been proposed to support multicore (Global Parallel Grub, Sequential Parallel Grub [], M-CACH[]).

# Global vs Partitioned ??

---

## Partitioned:

- Each core has its own ready-queue
- Single-core reservation & reclaiming strategy

## Pros

- Easy to implement: No migration
- a good exploitation of the platform

## Cons

- No-inter core reclaiming (under utilization of resources)

## Global:

- One ready Queue for all cores
- The  $m$  highest priority jobs are run

## Pros

- Implicit inter core reclaiming

## Cons

- Hard to implement,
- A big overhead due to the high number of migrations

# Problem formulation

---

## Our Solution

Is in the middle:

- Partition tasks to cores
- allow jobs to migrate to other cores to seek for extra-budget & tasks to migrate to have “a better” QoS.

We have:

- a set of tasks, each task is a set of jobs and :
  - A server, an abstract scheduling entity
  - The server contains necessary informations to schedule tasks:
    - Budget, Period, ...
- a set of identical cores

## Goal

Partition tasks and allow jobs to reclaim the unused bandwidth on all the cores (inter-core reclaiming).

# Plan

---

Introduction

Our Contributions

Results and discussions


Conclusion



## Single-Core Grub: Recall

Grub server uses 3 variables :

- Bandwidth  $u_i = \frac{Q_i}{p_i}$
- Server's deadline  $d_i$ . The task with the earliest  $d_i$  is executed.



- $V_i$  The virtual time, is updated when the task is in active-contending state:  
$$V_i(t + \Delta t) \leftarrow V_i(t) + \frac{U^a}{u_i} \Delta t, U^a = \sum_{S_i \in \{\text{ACT-CONT}\} \cup \{\text{ACT-N-CONT}\}} u_i.$$
- if  $V_i > d_i$ , the deadline is postponed as  $d_i = V_i + P_i(*)$ 
  - Task budget is exhausted.
  - Priority loss.

# Grub with Job migration

---

## Idea

- Before postponing the deadline, the job may seek to find extra-budget with the same urgency.
  - Keep the temporal isolation on the destination core
  - Migrate only when necessary (Sufficient budget is ensured in destination core)
- Each server is characterized by an extra variable :  $u_i^m$
- $u_i^m$  is the needed bandwidth to “eventually” complete before deadline on the destination core.

## When job migration is allowed

---

- A job is eligible for migration at time  $t_0$  if :

$$V_i(t_0) \geq d_i \quad (1)$$

$$d_i > t_0 \quad (2)$$

- Selection of destination core
  - The task must fit in destination core :  $u_i^m + U_{j'}^m(t) + U_{j'} \leq 1$
  - Ensure that with the current maximum load, the task can run *sufficient* execution time:  $u_i^m(d_i - t)/(u_i^m + U_m^a) > \epsilon$
- if this condition is not respected, migrating utilization can be reduced

# Job Migration

---

Before migration:

- New temporary server is created and initialized to `READY`
- Virtual time is set to  $t_0$ (migration moment)
- Server deadline is kept the same
- Migration flag is set to limit the migrations of the same job.

At the end of migration:

- the task returns immediately to its original processor at job completion.
- Migration flag is set to false

# Task migration

---

In the case of open systems:

- New tasks arrival  $\rightarrow$  load unbalance  $\rightarrow$  performance degradation

Solution

- *permanently* migrate the task to improve its quality of service.
- Permanent task migration can be performed at the end of the execution of its active job.
- A task permanent migration  $\rightarrow$  all its future jobs will start executing on the new processor.
- Permanent migration is triggered if  $miss\_rate^{now} > miss\_rate^{max}$

## Permanent migrations

---

Migrations are:

- immediate, if the migration core verifies:

$$u_i \leq 1 - U_j - U_j^m$$

- delayed, if previous condition not verified.
  1. Select a core that verifies:  $u_i \leq 1 - U_j$
  2. Lock migrations to the selected core
  3. Trigger the migration when the previous condition is verified
- aborted, if none of the previous conditions is verified.

# Plan

---

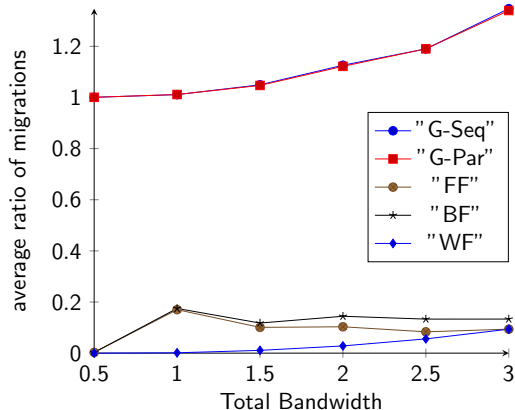
Introduction

Our Contributions

Results and discussions


Conclusion

## Job migration Vs Global scheduling: #Mig./job






## Job migration Vs Global scheduling: Dead-miss/job




## Job & task migration Vs Global scheduling: #Mis/job




## Job & task migration: #Mig./job

---




# On-line task arrival/leaving

## Active utilization



## Deadline miss ratio:



# Plan

---

Introduction

Our Contributions

Results and discussions

Conclusion

# Conclusion

---

- We proposed:
  - heuristics to partition a set of soft real-tasks to a set of identical cores
  - The execution time of each job is not known a priori
  - Tasks may enter and leave the system dynamically
- We allow:
  - Jobs to migrate to reclaim extra-budget
  - Tasks to migrate to

## Futur Work

---

- Implement the partitioned grub with job & task migration on Linux Kernel.
- Extend the current partitioned grub to hererogeneous architectures: ARM-bigLITTLE
- Complement the current version with DVFS and DPM to reduce the energy consumption.